

Class 5 - Sequential Regression

Continue the introduction of 1 new network/class

Sequential Regression

- Least Squares
- Logistic

Cost Functions

- Squared error
- Negative Log Likelihood

The homework will compare the behavior of sequential logistic regression using negative log likelihood (NLL) and squared error (SE) cost functions

Optimization

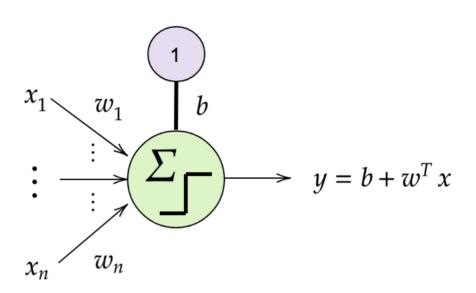
- Gradient descent
- Convexity

Perceptron

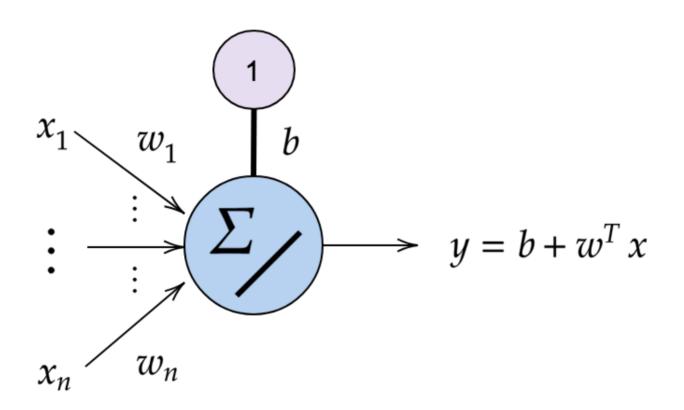
Last time we learned:

- Perceptron is trained using error correcting algorithm
- No gradient calculation
- Needs separable data

Today we'll look at 2 new networks

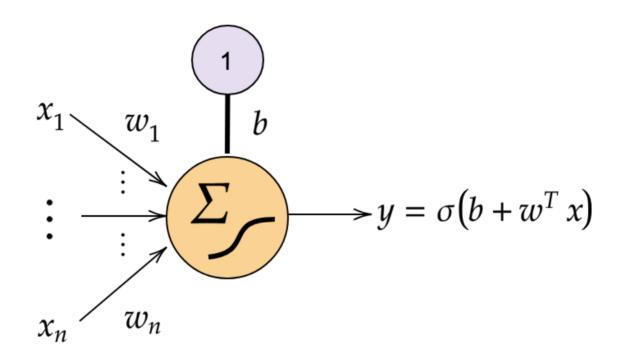


Least Squares Regression



$$egin{aligned} y^{(i)} &= b + \mathbf{w} \cdot \mathbf{x}^{(i)} & ext{Explicit Bias} \ y^{(i)} &= \mathbf{w}_0 + \mathbf{w} \cdot \mathbf{x}^{(i)} & ext{Implicit Bias} \ &= \mathbf{ ilde{w}} \cdot \mathbf{ ilde{x}}^{(i)} \end{aligned}$$

Logistic Regression



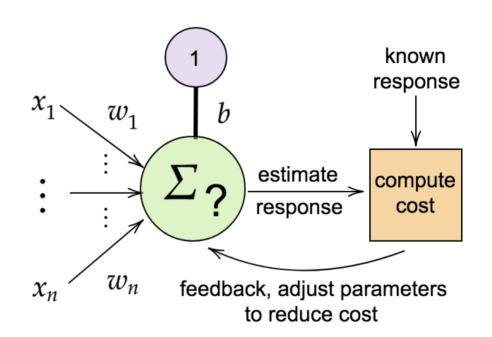
 $\sigma(\mathbf{x})$ is the sigmoid (logistic) function. Commonly used for the output layer for binary classification

 $\sigma(\mathbf{x}) \in [0,1]$ will be interpreted as a probability

Adjustable Parameters

Generic model with adjustable parameters

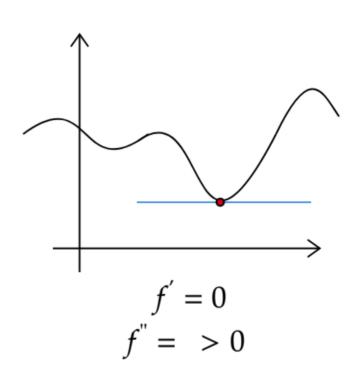
- The final output of the process is the model parameters
- Neural networks can have millions of parameters



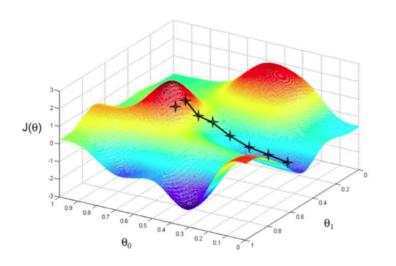
Optimization in 1-dimension

From 1^d calculus

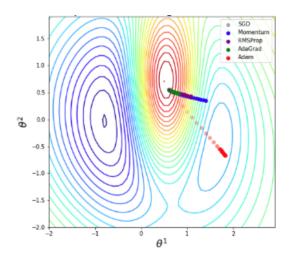
- If unconstrained, minimum satisfies f'(x) = 0 and $f''(\mathbf{x}) > 0$
- Points \mathbf{x} where $f'(\mathbf{x}) = 0$ are called stationary or critical points
- If constrained a minimum could also be on a constraint surface, like a boundary where $f' \neq 0$
- Networks have many parameters, so the cost function is minimized in a very high dimensional space.



A function of $f:\mathbb{R}^n o \mathbb{R}$ (scalar valued) defines a hypersurface in \mathbb{R}^n



Hypersurface defined by $f(\mathbf{x})$



Level curves of *f* define a topographical map

Directional Derivative

The directional derivative of $f: \mathbb{R}^n \to \mathbb{R}$ in the direction of a unit vector \mathbf{u} is defined as:

$$egin{aligned}
abla_u f &= \lim_{h o 0} rac{f(\mathbf{x} + h\mathbf{u}) - f(\mathbf{x})}{h} \ &= \lim_{h o 0} rac{f(\mathbf{x}_1 + h\mathbf{u}_1, \cdots, \mathbf{x}_n + h\mathbf{u}_n) - f(\mathbf{x})}{h} \end{aligned}$$

Gradient Vector

The gradient of f is a collection of directional derivatives. The components are equal to the directional derivatives in the coordinate basis directions:

$$rac{\partial f}{\partial \mathbf{x}_i} = \lim_{h o 0} rac{f(\mathbf{x} + h \hat{e}_i) - f(\mathbf{x})}{h}$$

Will show that derivatives in a direction ${\bf u}$ can be expressed in terms of the gradient

$$abla_{\mathbf{u}} f = \mathbf{u} \cdot
abla_{\mathbf{x}} f$$

Define $g(h) = f(\mathbf{x} + h\mathbf{u})$ for scalar h and unit vector \mathbf{u}

$$egin{aligned} g^{'}ig|_{0} &= \lim_{h o 0} rac{g(h) - g(0)}{h} \ &= \lim_{h o 0} rac{f(\mathbf{x} + h\mathbf{u}) - f(\mathbf{x})}{h} \ & riangleq \partial_{\mathbf{u}} f riangleq
abla_{\mathbf{u}} f \end{aligned}$$

Now use chain rule on g to set:

$$egin{aligned} g^{'}igg|_{h} &= (rac{\partial f}{\partial \mathbf{x}_{1}}\mathbf{u}_{1} + \cdots + rac{\partial f}{\partial \mathbf{x}_{n}}\mathbf{u}_{n})igg|_{\mathbf{x} + h\mathbf{u}} \ g^{'}igg|_{0} &= \mathbf{u} \cdot
abla_{f}igg|_{\mathbf{x}} \end{aligned}$$

$$abla_{\mathbf{u}} f = \mathbf{u} \cdot
abla f$$

The rate of change of the function f in the direction **u** is the scalar value $\nabla_{\mathbf{u}} f$

Want to show that the maximum value of $\|\nabla_{\mathbf{u}} f\|$ is in the direction of $\nabla_{\mathbf{x}} f$

$$egin{aligned}
abla_{\mathbf{u}}f &= \mathbf{u} \cdot
abla_{\mathbf{x}}f \\ |
abla_{\mathbf{u}}f| &= \|\mathbf{u}\| \|
abla f\| \cos \theta \\ &= \|
abla f\| \cos \theta \qquad \qquad \text{because} \|\mathbf{u}\| = 1 \end{aligned}$$

Where θ is the angle between \mathbf{u} and the gradient vector. It follows that $\|\nabla_{\mathbf{u}} f\|$ is maximized when $\cos(\theta) = 1$

$$egin{aligned} \max_{\|\mathbf{u}\|=1} \|
abla_{\mathbf{u}} f \| &= \|
abla_x f \| \ \mathbf{u} &= rac{
abla f}{\|
abla f \|} \end{aligned}$$

Gradient gives direction of most rapid increase.

Parametric Curves

Scalar valued functions are typically 'parameterized' by x. The graph of f can be written as:

$$(x,f(x))$$

$$\frac{d}{dx}(x,f(x))=(1,f')$$

$$(x_1,f(x_1))$$

$$(x_2,f(x_2))$$

$$1$$

So the vector described by the derivative of the parameterized expression is tangent to the curve provided that f' is defined.

Parametric Curves

Example

The equation of a circle can be parametrized as $(r\cos(\theta), r\sin(\theta))$. Want to show that the derivative of this $(-r\sin(\theta), r\cos(\theta))$ is tangent.

$$y=\sqrt{r^2-x^2}$$

$$y=\sqrt{r^2-x^2}$$
 $y'=rac{-x}{\sqrt{r^2-x^2}}$

Now use $x = r\cos(\theta), y = r\sin(\theta)$ to get

$$y' = -\frac{\cos \theta}{\sin(\theta)}$$

In general, if a curve in \mathbb{R}^n is described by a the parametric equation $\mathbf{x}(t)$, then \mathbf{x}' is tangent to $\mathbf{x}(t)$

Level Curves

A curve C in \mathbb{R}^n can be parametrized as $(\mathbf{x_1}(t), \dots, \mathbf{x_n}(t))$. The value of the function f along this curve is given by $f(\mathbf{x}(t))$ if the curve is a level curve of f then

$$f(\mathbf{x}(t)) = c$$

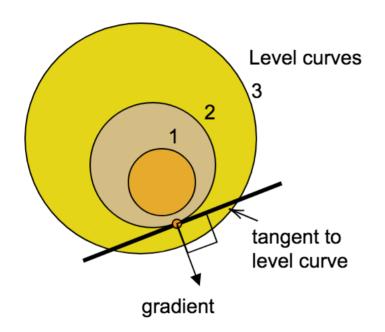
$$\frac{df}{dt} = \frac{\partial f}{\partial \mathbf{x_1}} \frac{d\mathbf{x_1}}{dt} + \dots + \frac{\partial f}{\partial \mathbf{x_n}} \frac{d\mathbf{x_n}}{dt} = \mathbf{x}' \cdot \nabla f$$

$$= \frac{dc}{dt} = 0$$

 $\mathbf{x}^{'}\cdot
abla f = 0$ shows that abla f is ot to the tangent vector $(\mathbf{x}^{'})$ to the level curve.

 ∇f is the direction of most rapid increase

 ∇f is perpendicular to level curves

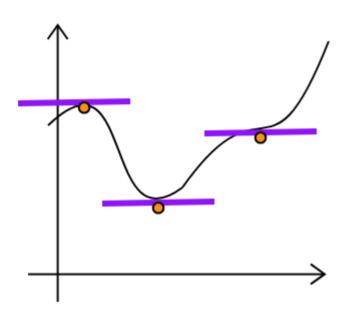


Gradient

- Points in direction of most rapid increase
- Is orthogonal to level curves

To move towards minimum, go in direction opposite to gradient.

Question: Is going down the steepest slope always the best choice?



From calculus

In 1^d stationary, or critical points, are defined as points c where f'(c)=0

A critical point is:

- A (local or global) minimum if f'' > 0
- A (local or global) maximum if f'' < 0
- An inflection or saddle point if f'' = 0

One-term Taylor approximation:

$$f(\mathbf{x}+\epsilon)pprox f(\mathbf{x})+\epsilon f^{'}(\mathbf{x})$$

It follows that a smaller value of f can be found if $\mathrm{sign}(\epsilon) = -\operatorname{sign}(f')$ and $f' \neq 0$

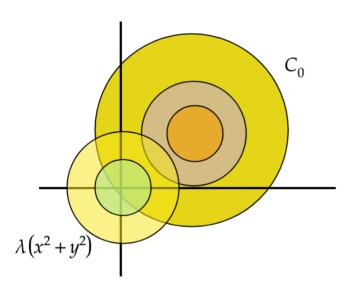
With each type of network, a scalar cost (objective) function C is defined. The goal will be to optimize cost by minimization.

- Neural Networks minimize a cost function. Minimization is just a convention.
 - \circ Minimizing C is the same as maximizing the negative of C.
- Minimization will be over a set of adjustable parameters.
 - ∘ *C* (**params**; data)
- When a cost function is linear in the adjustable **parameters**, can just solve f' = 0.

Gradient Descent

• For neural networks *C* will depend on nonlinear functions and have no explicit solution. Iterative solutions are needed.

Regularized Cost Function



For cost C, want to find point where $\nabla C = (\partial_x, \partial_y)C = 0$

Consider cost function C:

$$C_0=rac{1}{2}ig[(x-1)^2+(y-1)^2ig] \quad ext{base cost}$$

$$C = C_0 + \lambda (x^2 + y^2) \quad ext{cost plus penalty}$$

$$abla C = (x-1+2\lambda x\,,y-1+2\lambda y)$$

Assume the penalized minimum lies on the line x = y = t

$$C(t) = (t-1)^2 + 2\lambda t^2 = (2\lambda + 1)t^2 - 2t + 1$$
 $C' = 0 o t^* = rac{1}{(2\lambda + 1)}$ $abla C(t^*) = 0$

- The positive sum of convex functions is convex. Intuitive?
- By symmetry, the solution lies on the line connecting the individual minima

Multi-dimensional optimization

Neural network cost functions are defined over a high-dimensional parameter space. Critical points are defined using the gradient w.r.t the adjustable parameters.

Gradient over vector of parameters w:

$$\nabla_{\mathbf{w}}C = (\frac{\partial C}{\partial \mathbf{w}_1}, \cdots, \frac{\partial C}{\partial \mathbf{w}_n})$$

Critical point

$$abla_{\mathbf{w}}C = rac{\partial C}{\partial \mathbf{w}_i} = 0 \quad orall i$$

In general, will minimize C over adjustable parameters b and \mathbf{w}

Classification of critical points in higher dimensions

The Hessian is the matrix of second derivatives:

$$H =
abla_{\mathbf{w}}^2 C =$$

$$\begin{bmatrix} \frac{\partial^2 C}{\partial \mathbf{w}_1^2} & \cdots & \frac{\partial^2 C}{\partial \mathbf{w}_1 \partial \mathbf{w}_n} \\ \vdots & & \vdots \\ \frac{\partial^2 C}{\partial \mathbf{w}_n \partial \mathbf{w}_1} & \cdots & \frac{\partial^2 C}{\partial \mathbf{w}_n^2} \end{bmatrix}$$

$$= \begin{bmatrix} \frac{\partial (\nabla_w C)}{\partial \mathbf{w}_1} \\ \vdots \\ \frac{\partial (\nabla_w C)}{\partial \mathbf{w}_1} \end{bmatrix}$$

- Positive Definite $x^T H x > 0$
 - Local minimum -- Positive eigenvalues
- Negative Definite $\mathbf{x}^T H \mathbf{x} < 0$
 - Local maximum -- Negative eigenvalues
- Not Pos or Neg Definite
 - Saddle point -- Both positive and negative eigenvalues

Directional second derivative

$$egin{aligned} \partial_{\mathbf{u}}f &= \mathbf{u}\cdot
abla \mathbf{f} = (\mathbf{u}_1\partial_{x_1} + \ldots \mathbf{u}_n\partial_{x_n})f \ \partial_{\mathbf{u},\mathbf{u}}^2f &= \partial_{\mathbf{u}}(\mathbf{u}\cdot
abla f) = (\mathbf{u}_1\partial_{x_1} + \ldots \mathbf{u}_n\partial_{x_n})(\mathbf{u}_1\partial_{x_1} + \ldots \mathbf{u}_n\partial_{x_n})f \end{aligned}$$

Consider the i^{th} term

$$\mathbf{u}_i \partial_{\mathbf{x}_i} (\mathbf{u} \cdot
abla f) = \mathbf{u}_i (\mathbf{u}_1 \partial^2_{\mathbf{x}_i,\mathbf{x}_1} f + \dots + \mathbf{u}_n \partial^2_{\mathbf{x}_i,\mathbf{x}_n} f)$$

The term in parenthesis is the i-th row of the Hessian dotted with the direction vector \mathbf{u} . Since the total directional derivative is each of these terms weighted by \mathbf{u}_i we get:

$$\partial_{\mathbf{u},\mathbf{u}}^2 f = \mathbf{u}^T H \mathbf{u}$$
 where H is the Hessian of f

This form connects the multidimensional case with the one-dimensional case

Note: If the cost function depends on n adjustable parameters then H is $n \times n$ dimensional. For neural networks n can be extremely large (> 1,000,000).

Classification of critical points in n dimensions:

• Minima: H is positive definite

$$\mathbf{u}^T H \mathbf{u} > 0 \ \forall \mathbf{u} \neq 0$$

• Maxima: H is negative definite

$$\mathbf{u}^T H \mathbf{u} < 0 \ \forall \mathbf{u} \neq 0$$

• Saddle point: if H is neither positive definite or negative definite

Will cover positive definite matrices in future lecture. Useful facts:

- Positive definite matrices have positive eigenvalues
- Direction of maximum curvature will be the direction of the eigenvector of the largest eigenvalue

Newton's method successively minimizes a quadratic approximation to f given by 2 terms in a Taylor Series expansion

$$egin{aligned} f(\mathbf{x}) &pprox f(\mathbf{x}_k) +
abla f(\mathbf{x}_k)(\mathbf{x} - \mathbf{x}_k) + rac{1}{2}(\mathbf{x} - \mathbf{x}_k)^T H(\mathbf{x}_k)(\mathbf{x} - \mathbf{x}_k) \ d\mathbf{x} &= -H^{-1}(\mathbf{x}_k)(
abla f(\mathbf{x}_k))^T \end{aligned}$$

This is the multidimensional form of solving $y=\frac{1}{2}ax^2+b\mathbf{x}+c$ for y'=ax+b=0

$$egin{aligned} \mathbf{x}_{k+1} &= \mathbf{x}_k + d\mathbf{x} \ &= \mathbf{x}_k - H^{-1}(\mathbf{x}_k) (
abla f(\mathbf{x}_k))^T \end{aligned}$$

Newton's method converges much faster than steepest descent

- typically impractical for neural networks
- The Hessian can be an extremely large matrix
- Solving $H^{-1}(\mathbf{x}_k)(\nabla f(\mathbf{x}_k))^T$ is computationally expensive

Convexity

A **global minimum** is the point \mathbf{x}_0 such that

$$Cost(\mathbf{x}) \geq Cost(\mathbf{x}_0) \ \forall \mathbf{x}$$

A **local minimum** is a point \mathbf{x}_0 such that there is an $\epsilon > 0$, such that

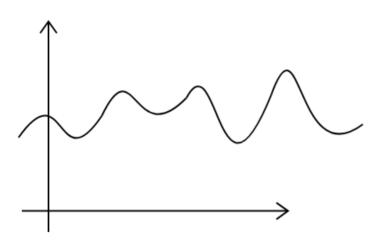
$$Cost(x) \geq Cost(x_0)$$
 when $|x - x_0| < \epsilon$

• Functions can have many local minima.

Convexity

- When the cost function is convex there are no local minima.
 - Any minimum is a global minimum.
- Will show later that
 - Linear regression with squared error cost is convex
 - Logistic regression with negative log likelihood cost is convex
- In general, Neural Network cost functions are not convex

Convexity



Lack of convexity for neural networks could be a deal breaker.

- In 1^d, finding global minimum of non-convex functions is NP-hard.
- For NN want to find the global minimum of the cost function in a high-dimensional parameter space
 - seems hopeless

It is not yet clear why neural networks work as well as they do.

A plausibility argument:

- Gradient descent gets stuck at local minima where $\nabla f = 0$
 - In 1-dimension, there's a 1 in 3 chance that a critical point is a local minima
 - In n dimensons there is only a 1 in 3^n chance that a critical point is a local minimum
 - In higher dimensions, it can be argued that almost all critical points are saddle points, so some component of the gradient will move the estimate away from the critical point.
 - Descent algorithms probably don't get stuck a saddle points
 - Convergence can be extremely slow

Convexity

A function $f: X \to \mathbb{R}^n$ is called convex if $\forall \mathbf{x}_1, \mathbf{x}_2 \in \mathbb{R}^n$, $t \in [0, 1]$

$$f(t\mathbf{x}_1+(1-t)\mathbf{x}_2) \leq tf(\mathbf{x}_1)+(1-t)f(\mathbf{x}_2)$$

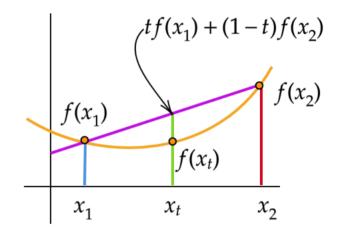
Local minima are global minima

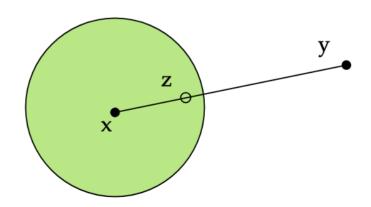
if \mathbf{x} is a local minimum then $f(\mathbf{z}) \geq f(\mathbf{x})$ for some neighborhood of \mathbf{x} . For any $\mathbf{y} \neq \mathbf{x}$ there is a $\lambda \in (0,1)$ such that $\mathbf{z} = \lambda \mathbf{y} + (1 - \lambda) \mathbf{x}$

$$egin{aligned} f(\mathbf{z}) &= f(\lambda \mathbf{y} + (1 - \lambda) \mathbf{x}) \ &\leq (1 - \lambda) f(\mathbf{x}) + \lambda f(\mathbf{y}) \ &= f(\mathbf{x}) + \lambda (f(\mathbf{y}) - f(\mathbf{x})) \end{aligned}$$

this implies

$$\lambda(f(\mathbf{y}) - f(\mathbf{x})) \geq 0$$





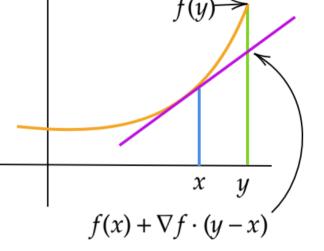
Convexity

Want to show that if f is convex and differentiable at x then

$$f(\mathbf{y}) \geq f(\mathbf{x}) +
abla f(\mathbf{x}) \cdot (\mathbf{y} - \mathbf{x})$$

Convexity condition

$$egin{aligned} f(lpha\mathbf{y}+(\mathbf{1}-lpha)\mathbf{x}) & \leq lpha f(\mathbf{y}) + (1-lpha)f(\mathbf{x}) \ f(\mathbf{x}+lpha(\mathbf{y}-\mathbf{x})) & \leq f(\mathbf{x}) + lpha(f(\mathbf{y})-f(\mathbf{x})) \ rac{f(\mathbf{x}+lpha(\mathbf{y}-\mathbf{x})) - f(\mathbf{x})}{lpha} & \leq f(\mathbf{y}) - f(\mathbf{x}) \end{aligned}$$



Example:

Left-hand side is derivative in direction y - x or

$$\nabla f(\mathbf{x}) \cdot (\mathbf{y} - \mathbf{x})$$

Maximum Likelihood (MLE)

Maximum Likelihood is one way to estimate model parameters from data.

- A parametric model choice $p(\mathbf{x}; \theta)$ is needed (Binomial, Gaussian, ...)
- An optimization process estimates the parameters θ given the data ${\bf x}$
- Has the intuitively appealing property that the estimated parameters maximize the likelihood of the observed data
- The assumption is that the data $\mathbf{x}^{(i)}$ are a representative sample from the distribution $p(\mathbf{x}; \theta)$
- MLE often has other desirable properties, such as being unbiased.

Maximum Likelihood

Parametric method that adjusts the parameters of an assumed model to maximize likelihood of the data. Given a set of data $\mathbf{x}^{(i)}$, assume the samples are drawn independently from a probability density $p(\mathbf{x}; \theta)$ with parameters θ .

The (\mathbf{x}^i) are independent, so their probability of occurring is

$$p(X; heta) = \prod_{i=1}^m p(\mathbf{x}^{(i)}; heta)$$

When considered as a function of θ instead of \mathbf{x} it is called the likelihood function.

• Likelihood is not a distribution over θ

Maximum Likelihood

Estimation

• Adjust θ to make the observed $\mathbf{x}^{(i)}$ most likely.

Log likelihood:

• Log is monotonic so can take the logarithm and not affect the location of the maximum.

$$egin{aligned} \mathcal{L}(heta;X) &= \ln l(heta;X) \ &= \sum_{i=1}^m \ln p(\mathbf{x}^{(i)}; heta) \end{aligned}$$

Minimization:

- The convention in neural networks is to minimize a cost function
 - So will maximize likelihood by minimizing negative log likelihood

Maximum Likelihood

Bernoulli

For
$$\mathbf{x}^{(i)} \in \{0,1\}$$

$$egin{aligned} p^{(i)} &= p(\mathbf{x}^{(i)}) = p^{\mathbf{x}^i} (1-p)^{1-\mathbf{x}^i} \ \mathcal{L}(p|\mathbf{x}) &= \ln\prod_{i=1}^M p^{\mathbf{x}^{(i)}} (1-p)^{1-\mathbf{x}^{(i)}} \ &= \sum_{i=1}^M \left(\mathbf{x}^{(i)} \ln p^{(i)} + (1-\mathbf{x}^{(i)}) \ln (1-p^{(i)})
ight) \ &= \sum_{i=1}^M \mathbf{x}^{(i)} \ln p + (M - \sum_{i=1}^M \mathbf{x}^{(i)}) \ln (1-p) \ &rac{d\mathcal{L}}{dp} = 0 \Rightarrow p = rac{1}{M} \sum_{i=1}^M \mathbf{x}^{(i)} \end{aligned}$$

Log Likelihood

- If $\mathbf{x}^{(i)}=1$ then i^{th} term is maximized by making p=1
- If $\mathbf{x}^{(i)} = 0$ then i^{th} term is maximized by making p = 0